

April 2025 Product Status

The product committee is responsible for improving PeeringDB's software as a whole, from bug fixes and small tweaks through the planning and execution of large changes to the web ui, the API, and peeringdb-py. The committee sits between users, who ask for improvements, and the developers, who implement them.

The committee over the past year or so has made an effort to recruit non-north-american members. Since not all members can make the monthly meeting, we've been trying to make asynchronous progress easier and do more business via the mailing list and github comments, with success. Our meetings have been more efficient as well, with less time spent rehashing the details and more time spent discussing implications and implementations.

We strongly encourage users to report their experience, good or bad, via github, email, or flagging down members at events. The more input we have, the better we can tailor results right out the gate, proactively, vs releasing imperfect behavior retroactively iterating. In 2024 we closed 144 issues, and are blazing through 2025 at a higher rate than ever following our nearly 50% increase in dev budget.

Over the past year we've made small and large improvements in several areas. A huge amount of effort went into fixing and improving search functionality- we think it's significantly better than it was, but please by all means send us a message if you get unexpected or missing results! Though the rate of errors is way down, search is not yet "feature complete".

Three quick search-related features worth mentioning: 1, you can choose in your user profile to not show any IXP results who lack presence in a facility for those of you with strong feelings about virtual IXs; 2, to be released very soon, you can get the curl command or python code to perform a particular hand-entered advanced search programmatically, via the API, copy-paste style. Finally 3, data quality issues are still top-of-rack, and we're continuing to sort out issues with naming similarities, street address and town name canonicalization, etc.

Another good mention is the recently-released user management API. Things like listing affiliated users, changing permissions, and adding and removing org user relationships are all now automatable via API. Again, please drop us a line if you are using this API heavily and either like it as-is or have feedback on improvements.

The last few accomplishments worth mentioning are peeringdb-py improvements, and the web ui's dark mode. Peeringdb-py has gotten more robust at exception handling and easier to install - we want to encourage as many of our users as feasible to use peeringdb-py for their API service vs PDB directly, since in most cases it's better for the user in addition to being cheaper/easier for us from the infra side. With respect to dark mode, check your user settings; you can choose to apply it either at night, or all the time.

Three important things coming down the pipe: 1, we are finally almost ready to release the phased-rollout web ui process, where 20% of users are opted into the beta ui at login but can choose whether to opt in all the time or never be chosen again; 2: API basic-auth is deprecated and will return errors starting this summer; and 3, multi factor auth will be required for all users starting July 1. Announcements will be sent in June for users that haven't yet successfully logged in via MFA, and some self service recovery flows will be provided, but please try to get your MFA configured and working well before then so you can tell us about any hiccups!

Huge thanks to all the PC members for keeping the train rolling; in no particular order, Terry Sweetser, Arnold Nipper, Matt Griswold, Jeff Bartig, Leo Vegoda, Martin Hannigan, Paul Hoogsteder, Peter Helmenstine, Steve McManus, Yan Berthier, and Yolandi Robinson.